

AURORA CRAMER

📍 Brooklyn, NY, United States of America

✉ aurora.linh.cramer@gmail.com

📷 auroracramer

🌐 auroracramer

👤 auroracramer.github.io

RESEARCH INTERESTS

Machine listening, music information retrieval, machine learning, digital signal processing

EDUCATION

New York University - New York, NY, USA

August 2017 - Present

PhD Candidate

GPA: 3.946

Department of Electrical and Computer Engineering

Advisor: Juan Pablo Bello

University of California, Berkeley - Berkeley, CA, USA

August 2011 - May 2015

Bachelor of Science (Honors)

GPA: 3.798

Department of Electrical Engineering and Computer Sciences

EECS Honors Program - Music/Audio

Advisor: David Wessel, Edmund Campion

RESEARCH EXPERIENCE

Music Audio Research Laboratory, NYU

September 2017 - Present

As a part of the machine listening team of the SONYC and BirdVox projects, investigating using natural structure in audio data from hierarchical relationships, task relationships, multi-modal dynamics, and embodied navigation to learn effective audio representations for machine listening tasks without large annotated datasets.

Center for New Music & Audio Technology, UC Berkeley

August 2014 - May 2015

Modeling musical sequences for the task of machine improvisation using an extension of author-topic modeling.

Video and Image Processing Lab, UC Berkeley

September 2013 - May 2014

Developing visualizations of indoor point cloud models and acquired sensor data for energy auditing applications

PEER-REVIEWED CONFERENCE AND WORKSHOP PUBLICATIONS

A. **Cramer**, M. Cartwright, F. Pishdadian, and J. P. Bello, "Weakly supervised source-specific sound level estimation in noisy soundscapes," in *2021 IEEE Workshop on Applications of Signal Processing to Audio and Acoustics (WASPAA)*, 2021.

M. Cartwright, A. **Cramer**, A. E. M. Mendez, Y. Wang, H.-H. Wu, V. Lostanlen, M. Fuentes, G. Dove, C. Mydlarz, J. Salamon, O. Nov, and J. P. Bello, "SONYC-UST-V2: an urban sound tagging dataset with spatiotemporal context," *Detection and Classification of Acoustic Scenes and Events 2020*, 2020.

A. **Cramer**, V. Lostanlen, A. Farnsworth, J. Salamon, and J. P. Bello, "Chirping up the right tree: incorporating biological taxonomies into deep bioacoustic classifiers," in *2020 IEEE International Conference on Acoustics, Speech and Signal Processing, ICASSP '20*, 2020.

M. Cartwright, A. E. M. Mendez, A. **Cramer**, V. Lostanlen, G. Dove, H.-H. Wu, J. Salamon, O. Nov, and J. P. Bello, "SONYC Urban Sound Tagging (SONYC-UST): a multilabel dataset from an urban acoustic sensor network," *Detection and Classification of Acoustic Scenes and Events 2019*, 2019.

V. Lostanlen, K. Palmer, E. Knight, C. Clark, H. Klinck, A. Farnsworth, T. Wong, A. **Cramer**, and J. P. Bello, "Long-distance detection of bioacoustic events with per-channel energy normalization," *Detection and Classification of Acoustic Scenes and Events 2019*, 2019.

M. Cartwright, A. **Cramer**, J. Salamon, and J. P. Bello, "TriCycle: audio representation learning from sensor network data using self-supervision," in *2019 IEEE Workshop on Applications of Signal Processing to Audio and Acoustics (WASPAA)*, 2019.

A. **Cramer**, H.-H. Wu, J. Salamon, and J. P. Bello, "Look, listen and learn more: design choices for deep audio embeddings," in *2019 IEEE International Conference on Acoustics, Speech and Signal Processing, ICASSP '19*, 2019.

C. Summers, G. Tronel, A. **Cramer**, A. Vartakavi, and P. Popp, “GNMID14: A collection of 110 million global music identification matches,” in *Proceedings of the 39th International ACM SIGIR Conference*, SIGIR ’16, 2016.

O. Oreifej, A. **Cramer**, and A. Zakhori, “Automatic generation of 3D thermal maps of building interiors,” in *ASHRAE*, 2014.

PATENTS

M. Cremer, A. **Cramer**, P. Popp, and C. Summers, “Responding to remote media classification queries using classifier models and context parameters,” July 6 2017. US Patent App. 15/185,616.

A. **Cramer**, M. Cremer, P. Popp, and C. Summers, “Model-based media classification service using sensed media noise characteristics,” July 6 2017. US Patent App. 15/185,654.

A. Vartakavi, C. Y. R. Gil, A. Gopakumar, and A. **Cramer**, “Methods and apparatus to generate recommendations based on attribute vectors,” May 27 2021. US Patent App. 16/695,169.

PROFESSIONAL EXPERIENCE

Mitsubishi Electric Research Laboratory - Cambridge, MA Speech & Audio Team - Audio Analysis and Source Separation Research Intern Investigating environmental source separation models.	August 2020 - January 2021
NVIDIA - Santa Clara, CA Applied Deep Learning Research - Research Intern Investigating audio inpainting methods using text-to-speech inspired deep sequence-to-sequence models.	May 2018 - August 2018
Gracenote - Emeryville, CA Applied Research - Audio Research Engineer Researching and developing audio classifiers to describe attributes of musical audio (e.g. genre classification, vocal detection, fingerprinting reliability); developing AWS applications for ingesting and processing audio content.	June 2015 - July 2017
Blue Jeans Network - Mountain View, CA Media Team - Media Software Engineering Intern Refactoring and improving the WebRTC and Speex noise suppression modules.	May 2014 - August 2014
Guidewire - Foster City, CA Development Operations - Software Engineering Intern Developing an optimization framework for managing virtual machines to balance cost and testing performance.	June 2013 - August 2013
WhereLab - Berkeley, CA Software Engineering Consultant Creating an interactive, wide-area augmented reality applications for iOS.	February 2013 - May 2013
Siemens Healthcare Diagnostics - Glasgow, DE Informatics Research and Development - Student Intern Creating a log parsing and statistical analysis application.	June 2012 - August 2012

HONORS AND AWARDS

ECE MS Student Award - New York University, Tandon School of Engineering	2018
Samuel Morse MS Fellowship - New York University, Tandon School of Engineering	2017
Music/Auto Challenge - Gracenote 5.0 Hackathon	2016
Auto Podcast Challenge - Gracenote 4.0 Hackathon	2015
3 rd Place - CSUA Hackathon, UC Berkeley	2013
3 rd Place - Code 4 Cal Hackathon, UC Berkeley	2013
Edward Frank Kraft Award - UC Berkeley	2011
Completed all A/B/C-Sides and Chapter 9 - Celeste (Nintendo Switch)	2022

SELECTED COURSE PROJECTS

Latent factor models for imputation of urban sound data	Fall 2019
Using Kalman filters and deep learning variants to model the temporal dynamics of audio embeddings computed from a large dataset of urban audio in order to impute embeddings for missing/corrupted audio.	
Ambisonic speech separation using recurrent neural networks using LSTMs	Spring 2019
Implementing an ambisonic speech separation method and trained and evaluated on synthesized ambisonic audio.	
Identifying and reducing gender bias in word-level language models	Spring 2018
Reducing bias in embeddings learned by a language model by applying regularization that penalizes projection onto an embedding subspace capturing variations in gender.	
Audio style transfer with cycle-consistent GANs	Spring 2018
Using a combination of WaveGAN and CycleGAN models for audio style transfer from raw audio.	
Online instrument source separation with source-filter models	Fall 2014
Developing an online framework for performing source separation of instruments in audio using source-filter models.	
Online instrument source separation with PLCA	Fall 2014
Developing an online framework for performing source separation of instruments in audio using PLCA.	

TEACHING EXPERIENCE

Teaching Assistant, ECE-GY 6143 Introduction to Machine Learning New York University	Fall 2018
Teaching Assistant, EE 126 Probability and Stochastic Processes University of California, Berkeley	Spring 2015
Teaching Assistant, EE 20N Structure and Interpretation of Signals and Systems University of California, Berkeley	Fall 2014

ACADEMIC SERVICE

Challenge Organization	
Task Organizer, IEEE AASP Challenge on Detection and Classification of Acoustic Scenes and Events	2020
Workshop Organization	
Student Volunteer, Workshop on Detection and Classification of Acoustic Scenes and Events	2019
Challenge Organization	
Task Organizer, IEEE AASP Challenge on Detection and Classification of Acoustic Scenes and Events	2019
Journal Reviewer	
IEEE Transactions on Audio, Speech and Language Processing	2019
Conference Reviewer	
Workshop on Detection and Classification of Acoustic Scenes and Events	2019
IEEE International Conference on Acoustics, Speech, and Signal Processing	2019

ORGANIZATIONS

IEEE Student Member	2019 - 2020
MIR @ Berkeley Cofounder University of California, Berkeley	2015
Computer Science Undergraduate Association Member University of California, Berkeley	2012 - 2015
Eta Kappa Nu Honor Society Member University of California, Berkeley	2012 - 2015

SELECTED OPEN SOURCE PROJECTS

BirdVoxClassify

Open-source Python library for performing classification of avian flight calls at different taxonomic levels

SONYC Urban Sound Tagging Dataset (SONYC-UST)

Urban sound classification dataset, released as part of the DCASE 2019 Challenge Task 5: Urban Sound Tagging and DCASE 2020 Challenge Task 5: Urban Sound Tagging with Spatiotemporal Context.

openl3

Open-source Python library for extracting audio and image embeddings, using pre-trained models based on the Look, Listen, and Learn approach

SELECTED COURSEWORK TOPICS

Machine Learning & AI	Statistical Signal Processing	Music Perception and Cognition
Machine Listening & MIR	Linear Dynamical Systems	Computer Music
3D Audio	Time Series Analysis	Compilers and Languages
Digital Signal Processing	Statistical Learning Theory	Parallel Programming
Probability & Stochastic Processes	Data Structures & Algorithms	Scientific Writing

PROGRAMMING AND DEVELOPMENT SKILLS

Python	Web (HTML, CSS, JavaScript, Flask)
MATLAB	Data visualization (matplotlib, d3)
C/C++	UNIX scripting
Java	AWS (ElasticBeanstalk, S3, DynamoDB, CloudWatch)